

Regular Expressions and Finite Automata

CSCI 3136: Principles of Programming Languages

Agenda

- Announcements
- Readings:
 - Today: 2.2.1 (From optional text)
 - Next: 2.2.1
 - Note: I recommend using alternative texts for this part of the course:
 - E.g., Hopcroft et al, “Introduction to Automata Theory”
- Lecture Contents
 - Motivation
 - Regular Expressions
 - Deterministic Finite Automata
 - Nondeterministic Finite Automata

Specifying Regular Languages

- There are many ways to specify a regular language
 - $\{a, ab, abc\}$
 - $\{a\}^*$
 - $\{a, b, c\}^*$
 - $\{1^n 0 \mid n > 0\}$
 - Set of all positive integers (base 10)
 - $\{\Sigma^* @ \Sigma^* (. \Sigma^*)^n \mid \Sigma = \{a, b, c, \dots, z\}, n \geq 0\}$
- Problems:
 - The specification is not standard
 - Hard for a program to interpret the specifications above.
- Does some of the notation above look familiar?

Regular Expressions

- Idea: Regular expressions (REs) are a concise way to specify regular languages
- **Theorem:** L is regular if and only if there is a regular expression R that specifies L
- Recursive definition:
 - Base cases:**
 - \emptyset defines the empty language (no words)
 - $a, a \in \Sigma$ defines the language $\{a\}$
 - ϵ defines the language $\{\epsilon\}$
 - Inductive step:** If R, R_1 , and R_2 , are REs:
 - $R_1 | R_2$ defines $L_1 \cup L_2$, (union) where R_i specifies L_i
 - $R_1 R_2$ defines $L_1 L_2$, (concatenation) where R_i specifies L_i
 - R^* defines L^* , (Kleene-*) where R specifies L

Examples of Regular Expressions

| | <u>Corresponding Language</u> |
|--|------------------------------------|
| • $ab c$ | $\{ab, c\}$ |
| • $a(b c)$ | $\{ab, ac\}$ |
| • a^* | $\{a\}^*$ |
| • $(a b c)^*$ | $\{a,b,c\}^*$ |
| • 11^*0 | $\{1^n0 n>0\}$ |
| • $0^*[1-9][0-9]^*$ | Set of all positive integers |
| • $aa^*bbb^*cccc^*$ | $\{a^ib^jc^k i>0, j>1, k>2\}$ |
| • $0^*(100^*)^*(1 \epsilon)$ | Binary strings with no adjacent 1s |
| • $[a-z][a-z]^*@[a-z][a-z]^*(.[a-z][a-z]^*)^*$ | email address |

Note: Notation $[a-z] = (a|b|c|d|\dots|z)$

Applications and History


- Applications:
 - Search (and replace)
 - editors, string manipulation libraries,
 - scanners
 - specification of tokens.
- History
 - Stephen Cole Kleene, 1956
 - “Representation of events in nerve nets and finite automata”
 - Ken Thompson developed editors: QUE, ed, grep
 - Used in awk, emacs, vi, lex, etc...
 - Henry Spencer, 1986, C regex library used in Tcl, Perl, etc...

How to Build a Scanner

- We now have a standard (machine-friendly) way to specify regular languages.
- So what?
- We now need a way to decide if a given string σ is in a given regular language L .
- How do we do this?
- We use a *Deterministic Finite Automata* (DFA).

Deterministic Finite Automata


- A DFA M is a machine that
 - Takes a string $\sigma \in \Sigma^*$ as input
 - Either accepts σ if $\sigma \in L$
 - Or rejects σ if $\sigma \notin L$
- M recognizes L if it accepts σ if and only if $\sigma \in L$
- A DFA consists of:
 - set of states
 - *start* state
 - set of *final* states
 - transition function




Operation of a DFA

A DFA

- Starts in the start *state*
- Reads in string σ one character at a time
- Computes the next state based on current state and character
- *Transitions* to the next state
- *Accepts* σ if and only if it is in a *final* state after reading σ .



What Language Does this DFA Recognize?



$(0|1)^*1$

Formal Definition of a DFA

- A DFA is a 5-tuple $M = (Q, \Sigma, \delta, q_0, F)$
 - Q set of states
 - Σ alphabet
 - δ transition function (complete): $\delta : Q \times \Sigma \rightarrow Q$
 - q_0 start state, $q_0 \in Q$
 - F set of final states, $F \subseteq Q$
- A DFA M accepts a string $\sigma \in \Sigma^*$ if and only if it is in a final state after reading σ .
- A DFA M recognizes language L if and only if it only accepts all the strings in L
$$L(M) = \{\sigma \in \Sigma^* \mid M \text{ accepts } \sigma\}$$

Examples of DFAs

DFA that accepts all binary strings that have no consecutive 1s.

$M = (\Sigma, Q, \delta, q_0, F)$

- $\Sigma = \{0, 1\}$
- $Q = \{q_0, q_1, q_2\}$
- $F = \{q_0, q_1\}$

• δ :

| State | Symbol | New State |
|-------|--------|-----------|
| q_0 | 0 | q_0 |
| q_0 | 1 | q_1 |
| q_1 | 0 | q_0 |
| q_1 | 1 | q_2 |
| q_2 | 0 | q_2 |
| q_2 | 1 | q_2 |

DFA that accepts $L = (0|1)^*1$

$M = (\Sigma, Q, \delta, q_0, F)$

- $\Sigma = \{0, 1\}$
- $Q = \{q_0, q_1\}$
- $F = \{q_1\}$

• δ :

| State | Symbol | New State |
|-------|--------|-----------|
| q_0 | 0 | q_0 |
| q_0 | 1 | q_1 |
| q_1 | 0 | q_0 |
| q_1 | 1 | q_1 |

More Examples

- $L \subseteq \{a, b\}^*$: all strings containing an odd number of b's
- $L \subseteq [0 - 9]^*$: all integers divisible by 100
- $L \subseteq [a - z, @.]^*$: all valid email addresses
- $L \subseteq \{0, 1\}^*$: all binary numbers not divisible by 3

Nondeterministic Finite Automata (NFA)

- A DFA is deterministic in that it has a single transition for each symbol and state
I.e., A DFA traces a single path for each input
- An NFA is like a DFA except it may have a choice of transitions for a given state and character.
I.e., An NFA may trace multiple paths for an input
- Two kinds of nondeterministic choices:
 - **ϵ transitions:** transition to another state without reading a character
 - **multiple successor states:** multiple transitions to different states from same state and same character
- An NFA *accepts* a string σ if one of the paths ends in a final state

Example of an NFA


NFA that accepts binary strings ending in 01 or 001

$$L = (0|1)^* 0(1|01)$$

$$M = (\Sigma, Q, \delta, q_0, F)$$

- $\Sigma = \{0,1\}$
- $Q = \{q_0, q_1, q_2, q_3\}$
- $F = \{q_3\}$
- δ :

| State | Symbol | New State |
|-------|------------|-----------|
| q_0 | 0 | q_0 |
| q_0 | 0 | q_1 |
| q_0 | 1 | q_0 |
| q_1 | 0 | q_2 |
| q_1 | ϵ | q_2 |
| q_2 | 1 | q_3 |



Formal Definition of an NFA

- An NFA is a 5-tuple $M = (Q, \Sigma, \delta, q_0, F)$
 - Q set of states
 - Σ alphabet
 - δ transition function: $\delta : Q \times (\Sigma \cup \{\epsilon\}) \rightarrow 2^Q$
 - q_0 start state, $q_0 \in Q$
 - F set of final states, $F \subseteq Q$
- Every input σ induces a set of paths traced by δ as σ is read
- NFA M accepts a σ if and only if one of the paths ends in a final state
- NFA M *recognizes* $L(M) = \{\sigma \in \Sigma^* \mid M \text{ accepts } \sigma\}$
- Question: Are NFAs more powerful than DFAs?


NFA Example 1

NFA that accepts $L = (0|1)^*1$

$M = (\Sigma, Q, \delta, q_0, F)$

- $\Sigma = \{0,1\}$
- $Q = \{q_0, q_1\}$
- $F = \{q_1\}$
- δ :

| State | Symbol | New State |
|-------|--------|-----------|
| q_0 | 0 | q_0 |
| q_0 | 1 | q_0 |
| q_0 | 1 | q_1 |




NFA Example 2

NFA that accepts $L = (0|1)^*(010|101)$

$M = (\Sigma, Q, \delta, q_0, F)$

- $\Sigma = \{0,1\}$
- $Q = \{q_0, q_1, q_2, q_3, q_4, q_5\}$
- $F = \{q_5\}$
- δ :

| State | Symbol | New State |
|-------|--------|-----------|
| q_0 | 0 | q_0 |
| q_0 | 1 | q_0 |
| q_0 | 0 | q_1 |
| q_0 | 1 | q_3 |
| q_1 | 1 | q_2 |
| q_2 | 0 | q_5 |
| q_3 | 0 | q_4 |
| q_4 | 1 | q_5 |



More NFA Examples

- $L \subset \{a, b\}^*$: all strings containing an odd number of a's or b's
- $L \subset \{0, 1\}^*$: all binary strings containing the substring 101 or 010

Are these all the same?

- We have discussed a variety of specifications: RLs, RE, DFAs, NFAs
 - RLs: a class of languages
 - RE a way to specify RLs
 - DFAs: a way to implement scanners for RLs
 - NFAs: a simpler way to implement scanners for RLs
- Questions:
 - Are these all of equal power?
 - Are NFAs same as DFAs?
 - Do REs specify only regular languages?

Regular Languages Equivalence Theorem

- Thm: The following statements are equivalent:
 - i. L is a regular language.
 - ii. L is the language described by a regular expression.
 - iii. L is recognized by an NFA.
 - iv. L is recognized by a DFA.
- We will prove: $(i) \equiv (ii) \equiv (iii) \equiv (iv)$

Regular Languages are equivalent to Regular Expressions

- Every regular language can be specified by a regular expression.
- Every regular expression specifies a regular language.
- Idea: There is a one-to-one correspondence between definitions of RLs and REs
- Apart from notation, the recursive definitions are identical.

| Operation | Regular Language | Regular Expression |
|------------------|-----------------------|--------------------|
| Empty Language | \emptyset | \emptyset |
| Empty String | $\{\epsilon\}$ | ϵ |
| Single character | $\{a\}, a \in \Sigma$ | a |
| Disjunction | $L_1 \cup L_2$ | $R_1 R_2$ |
| Concatenation | $L_1 L_2$ | $R_1 R_2$ |
| Kleene-* | L^* | R^* |