

# Exceptions and Continuations

CSCI 3136: Principles of Programming Languages

# Agenda

- Announcements
  - **SRI's Today!!!**
  - Assignment 9 is due July 30
  - **Final exam, 1:00pm, Friday, August 2 in CHEB 170**
- Readings: Read Chapter 6.6, 9
- Lecture Contents
  - Exceptions
  - Motivation for Continuations
  - Continuations
  - Continuations in Scheme
  - Uses and Abuses of Continuations
  - Implementation
  - Co-routines

# How are the Student Ratings of Instruction (SRI) used?

- ✓ Course and program ***(re) design***.
- ✓ ***Evaluation*** of teaching effectiveness.
- ✓ ***Promotion and tenure applications*** for instructors, and other personnel decisions.
- ✓ Preparation of supporting evidence for ***teaching awards and grants***.
- ✓ ***Quality assurance*** processes in the review and restructure of institutional, faculty, department and program goals.

# How to complete the SRI

- Find the email in your Dal email account
  - Subject heading (depending on the system) is:
    - *Student Ratings of Instruction; or*
    - *Course Name and Number*
- Open the email and click on the link
  - Your course list should be visible
- Select the course for which you want to complete the evaluation
- Be sure to hit the SUBMIT button when you FINISH completing the form
- You may also SAVE and return to your work later

Also available via Brightspace

# Return Values

- In most languages functions typically return r-values
  - A value that can be assigned to a variable or used in an expression
- Some languages, such as C++, allow functions to return l-values (locations of the value)
  - Seen in a previous lecture
- Return of l-values can be simulated in most languages
  - Using pointers in C
  - Returning references in Java
  - Etc.
- But ... Sometimes it's hard to know what to return!

# Exception Handling

- Things go wrong (bleep happens), we need to handle it gracefully
- *Exception* are unexpected or abnormal conditions during execution
  - Generated automatically in response to runtime errors
  - Raised explicitly by the program
- Exception handling is needed to
  - Perform operations necessary to recover from the exception
  - Terminate the program gracefully
  - Clean up resources allocated in the protected block
- Exception handling allows the programmer to
  - Specify what to do when an error occurs during program run-time
  - Separate the common path code from the error handling code

# Exception Handling Syntax

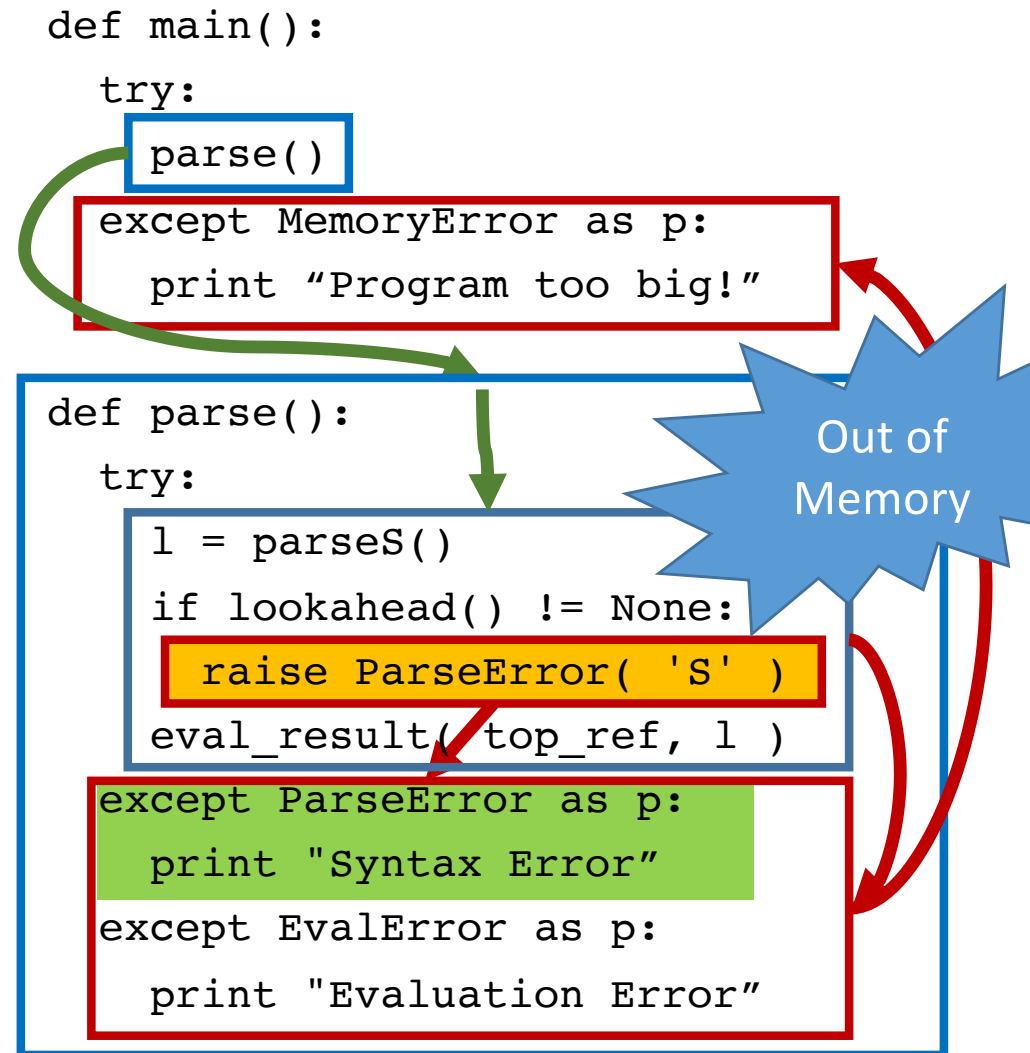
- Syntax for catching and handling exceptions tends to be similar
- A **protected block** comprises 3 parts:
  - **try** : the common path code to be executed
  - **catch** : exception handlers for each exception to be caught
  - **finally** : an optional "clean-up" handler that always runs after the "try" regardless of whether an exception occurs
- Exception are **raised** (or thrown) by a raise (or throw) statement

**raise Exception\_1(...)**

```
try {  
    // common path  
} catch ( Exception_1 e ) {  
    // Exception 1 handler  
} catch ( Exception_2 e ) {  
    // Exception 2 handler  
} ...  
} else { // optional  
    // default handler  
} finally { // optional  
    // clean up code  
}
```

# Exception Handling Semantics

- An exception handler is lexically bound to a block of code
- When an exception is raised in the block, search for a handler in present scope
- If there is no matching handler in present scope,
  - The scope is exited (may include block or subroutine)
  - A handler is searched for in the next scope





# Language Support

- How are exceptions represented?
  - Built-in exception type (Python)
  - Object derived from an exception class (Java)
  - Any kind of data can be passed as part of an exception
- When are exceptions raised?
  - Automatically by the run-time system as a result of an abnormal condition
    - e.g., division by zero, null dereference, out-of-bounds, etc
  - `throw` or `raise` statement to raise exceptions manually
- Where can exceptions be handled?
  - Most languages allow exceptions to be handled locally
  - Propagate unhandled exceptions up the dynamic chain.
    - Clu does not allow exceptions to be handled locally
- Some languages require exceptions that are thrown but not handled inside a subroutine be declared (Java)

# Language Non-support

- Some languages do not support exceptions  
e.g., C
- Solution 1:
  - Reserve a return value to indicate an exception
- Solution 2:
  - Caller passes a closure (exception handler) to call
- Solution 3:
  - In C, signals and `setjmp` / `longjmp` can be used to simulate exceptions

```
#include <setjmp.h>
```


```
int func(...) {  
    static jmp_buf env;  
    int i = setjmp(env);  
    if( i == 0 ) {  
        /* common path */  
        ...  
        /* exception 42 */  
        longjmp(env, 42);  
        ...  
    } else if( i == 42 ) {  
        /* handle exception 42 */  
    } else if( i == ? ) {  
        ...  
    }  
}
```

```
graph TD  
    A["int i = setjmp(env);"] -- green arrow --> B["longjmp(env, 42);"]  
    B -- red arrow --> A  
    B -- blue arrow --> C["} else if( i == 42 ) {  
    /* handle exception 42 */  
}"]
```

# Exception Implementations

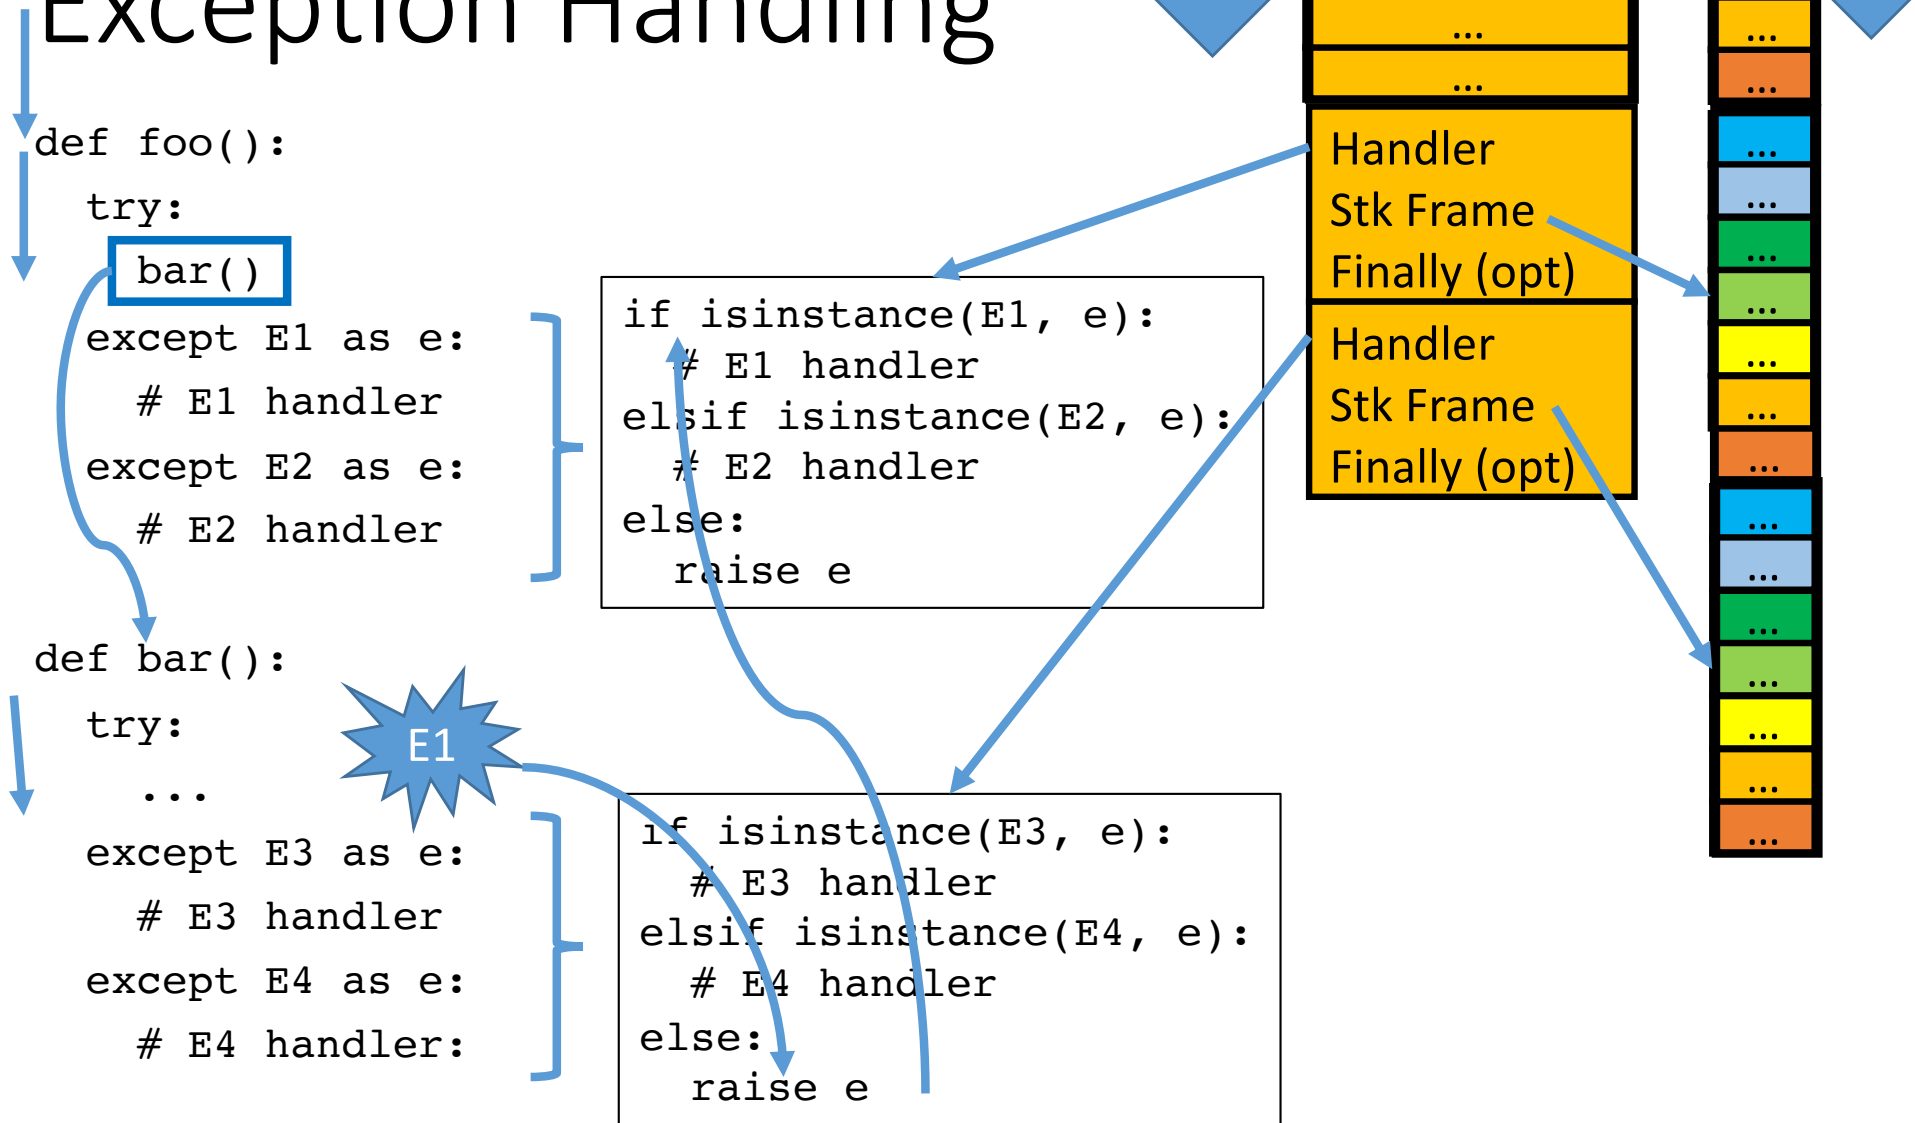
- Options:
  - Simple (Pay as you go)
  - Location to Exception map (Pay on Exception)
  - Hybrid

# Simple, Pay-as-You-Go Exception Handling



- Idea:
  - The program uses a second stack, called a Handler Stack (HS)
  - When a protected block is entered, a handler is pushed on the (HS)
    - Pointer to the handler code
    - Current stack frame (Program Stack)  
I.e., referencing environment
  - Sound familiar?
    - An optional exit (finally) handler may also be pushed
  - If there are multiple exception handlers, these are implemented using an if/elseif/... construct in a single handler
  - When a protect block is exited, the handler is popped of the stack
- Simple implementation is costly because handler stack is manipulated on entry/exit of each protected block

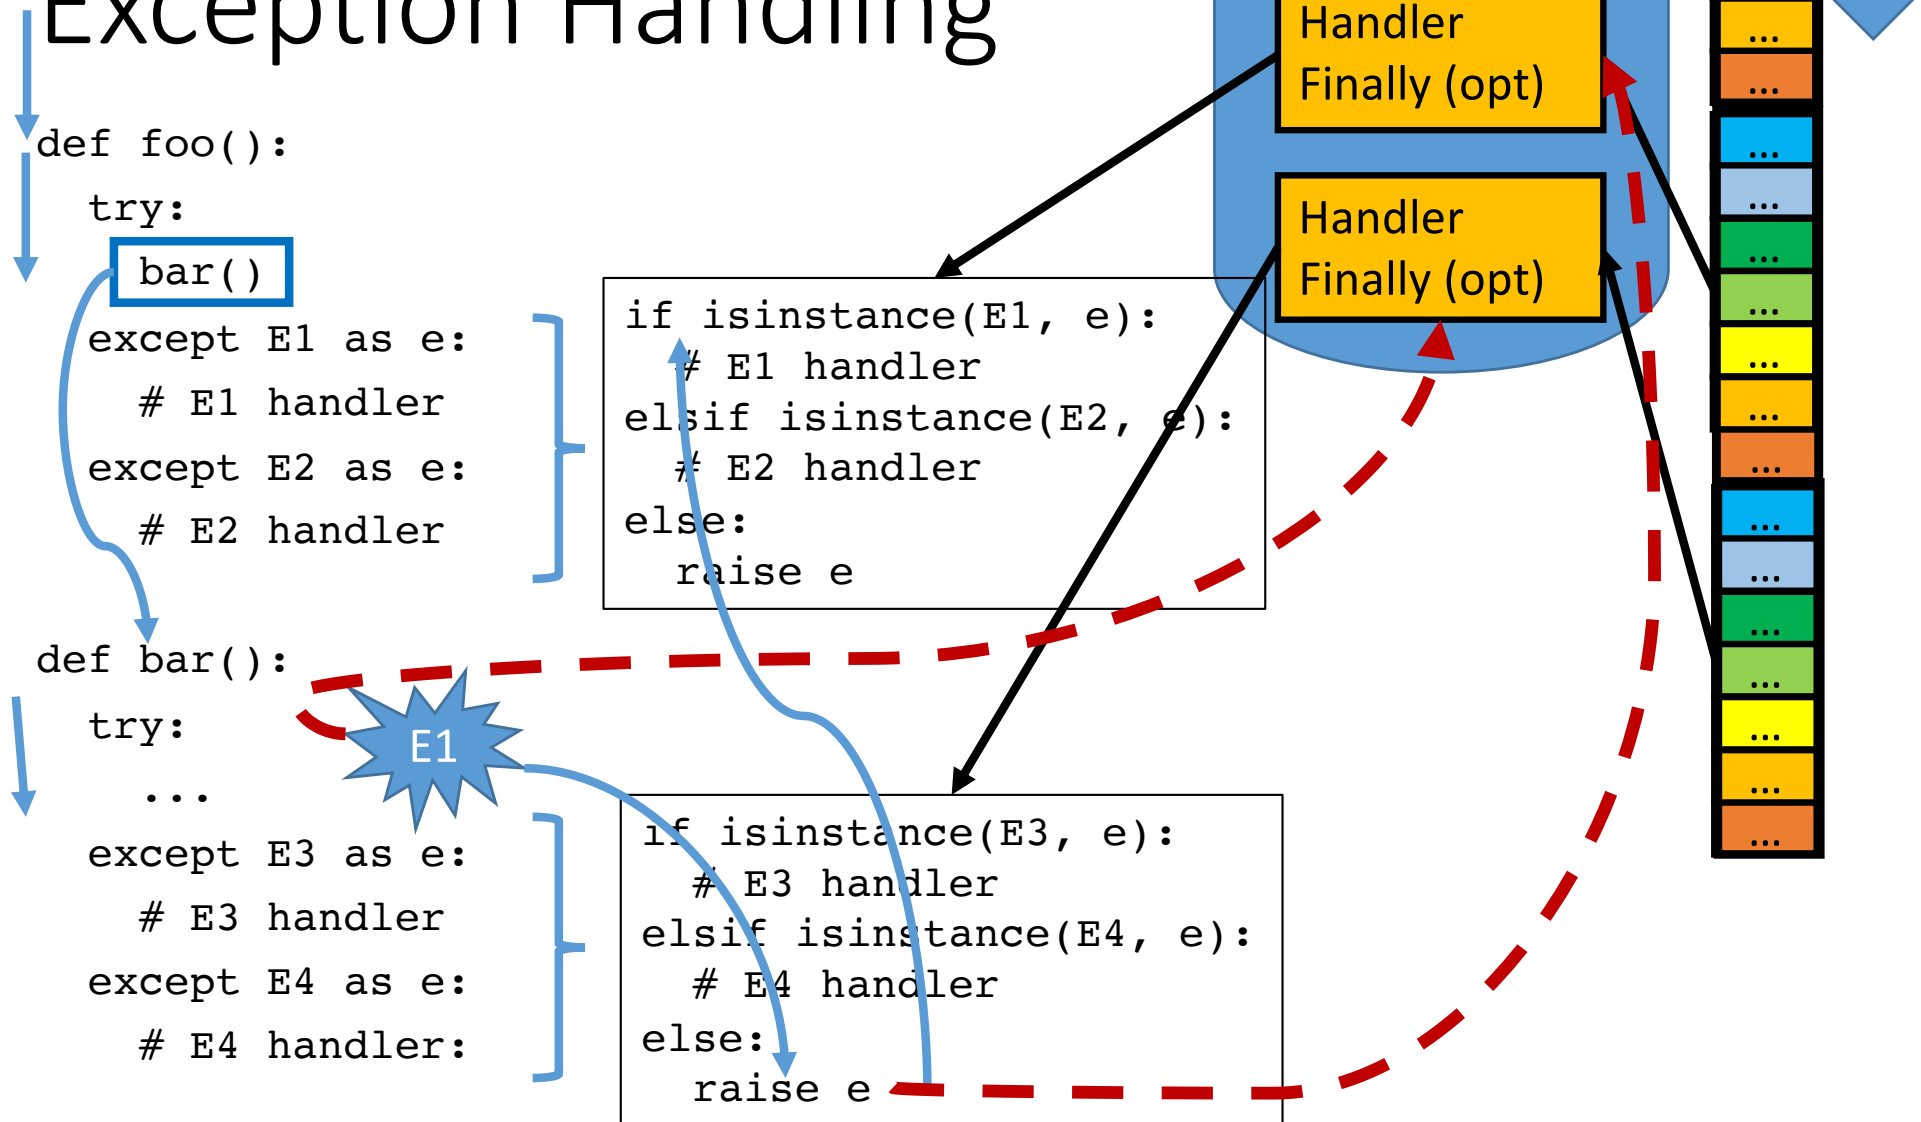
# Simple, Pay-as-You-Go Exception Handling



# Location to Exception Mapping

- A faster implementation (Pay on exception)
- Store a global map of code blocks (memory addresses) to handlers
  - Generated by compiler/linker
- On exception, index map with program counter to get handler
- Still need to keep track of stack frames
  - Each stack frame stores a pointer to most recent protected block

# Pay on Exception Exception Handling



# Comparison of the 2 Approaches

- Location-based Exception handling
  - Handling an exception is more costly (search), but exceptions rare
  - No cost if no exceptions
  - Cannot be used if the program consists of separately compiled units and the linker is not aware of this exception handling mechanism
- Hybrid Approach:
  - Use a local map for each subroutine
  - Store a pointer to a local map in subroutines stack frame





# Motivation


- Its useful to have a general way to implement a variety of mechanisms such as
  - Exceptions
  - Gotos
  - Coroutines
  - Subroutines
  - Closures
  - Transform recursion into tail recursion
  - etc
- Continuations are a general mechanism for doing this.

# Continuation

- A continuation is the “future” of the current computation
- Represented as the current
  - Stack contents (sequence of stack frames)
  - Referencing environment
  - Current program state
    - Program counter (current location)
    - Registers
    - Etc.
- Analogy: A bookmark in the computation
  - A way for a program to return to the location in your program, as if nothing has happened
- Analogy: A “back” button for a program



# We Already Use Bookmarks!





# A Picture of a Continuation

```
( define find_neg
  ( lambda ( L )
    ( define finder ( lambda ( exit )
      ( define do_check ( x )
        ( if ( negative? x)
              ( exit x) ) )
      ( for-each do_check L )
      #t )
    ( call/cc finder ) ) )
```

## Continuation:

- Program State
  - Location
  - Registers
- Referencing Environment
- Stack Frames



# Aside: Current Program State



- What does the immediate behaviour of a program depend on?
  - l.e., what affects
    - Which instruction is executed next?
    - The result of the instruction?
- Answer: the CPU State
  - Program counter (location of next instruction)
  - General purpose registers (current values being manipulated by the program)
  - Stack pointer register (current stack frame)
- Idea: Current program state corresponds to the current state of the CPU, including current location in the program

# A Picture of a Continuation

```
( define find_neg
  ( lambda ( L )
    ( define finder ( lambda ( exit )
      ( define do_check ( lambda ( x )
        ( if ( negative? x)
              ( exit x) ) ) )
      ( for-each do_check L ) a
      #t )
    ( call/cc finder ) ) )
```

## Continuation:

- Program State
  - Location
  - Registers
- Referencing Environment
- Stack Frames



# Continuation as First-Class Objects

- A *first-class* object is a something that can be passed to a function and returned by a function
- Continuations are first-class objects in Scheme:
  - Passed as function arguments
  - Returned as function results
  - Stored in variables and data structures
- Note: A continuation can be “resumed” from *anywhere* in the program
  - Just like flipping to a bookmark can be done from anywhere in a book!

# Continuations in Scheme

- Continuations are created by
  - Taking a snapshot of the current state of the program  
i.e., creating the continuation
  - Calling a function and passing it the snapshot
- In Scheme this is done with the special function:

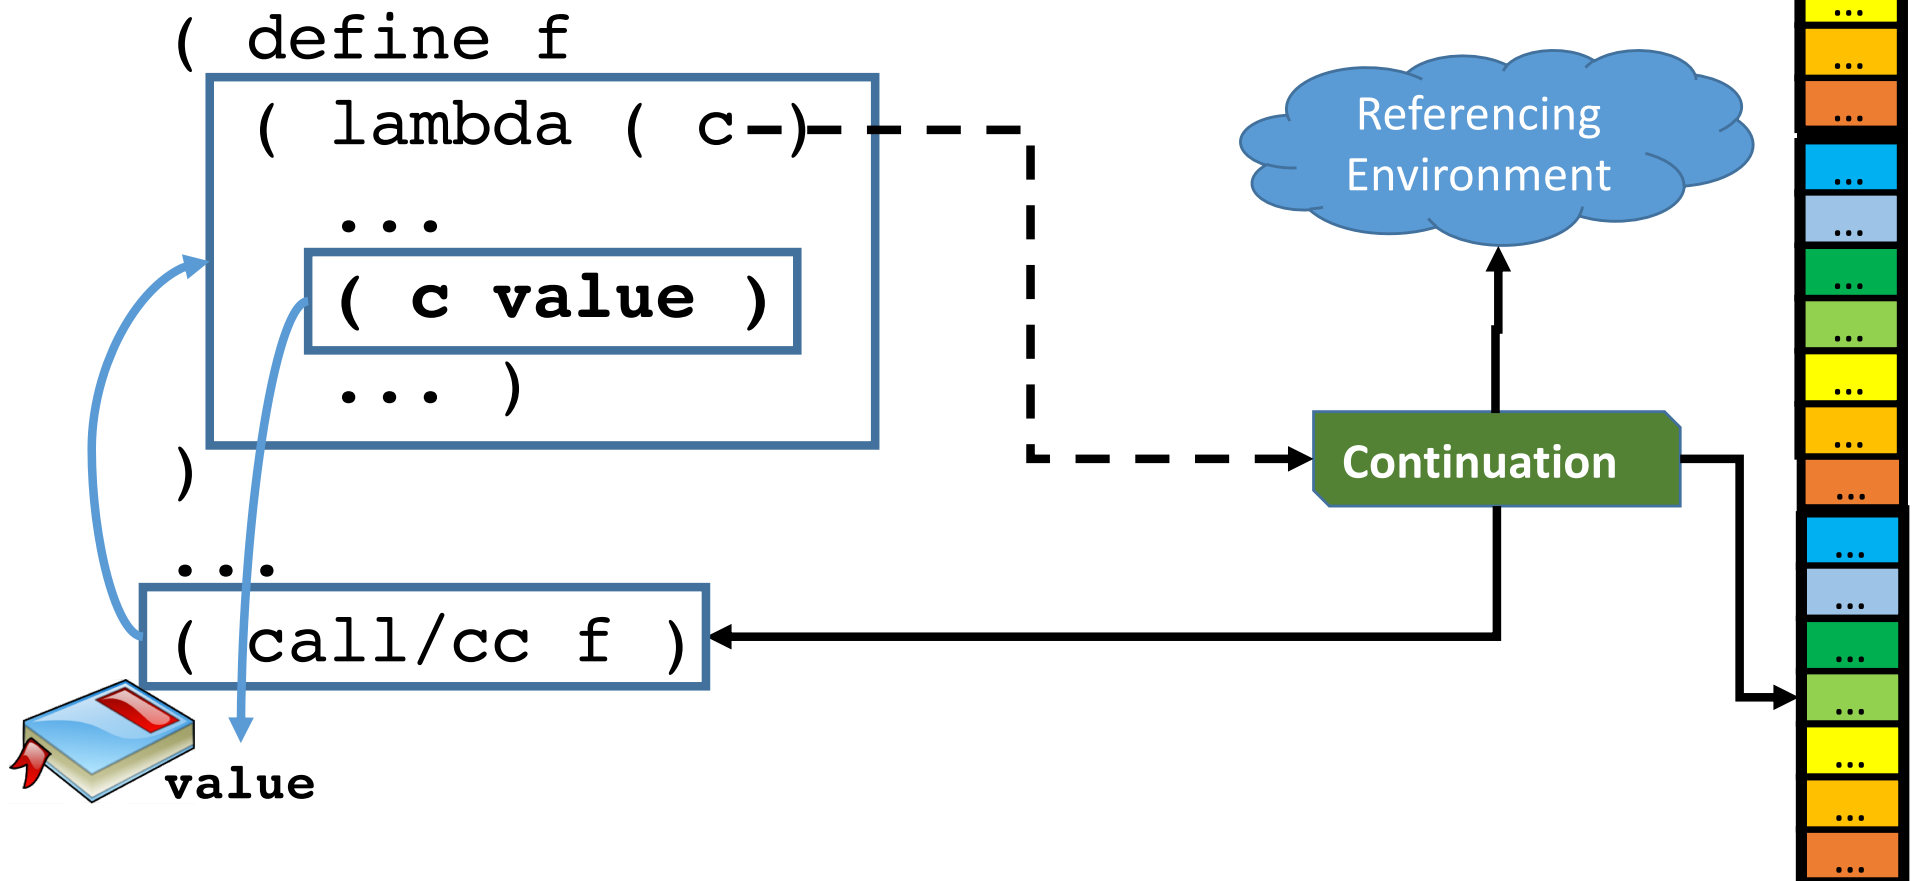
```
( call-with-current-continuation f )
```

- Calls function  $f$
- Passes the current continuation to  $f$  as an argument

Short form: ( call/cc f )



# What Does `(call/cc f)` Do?

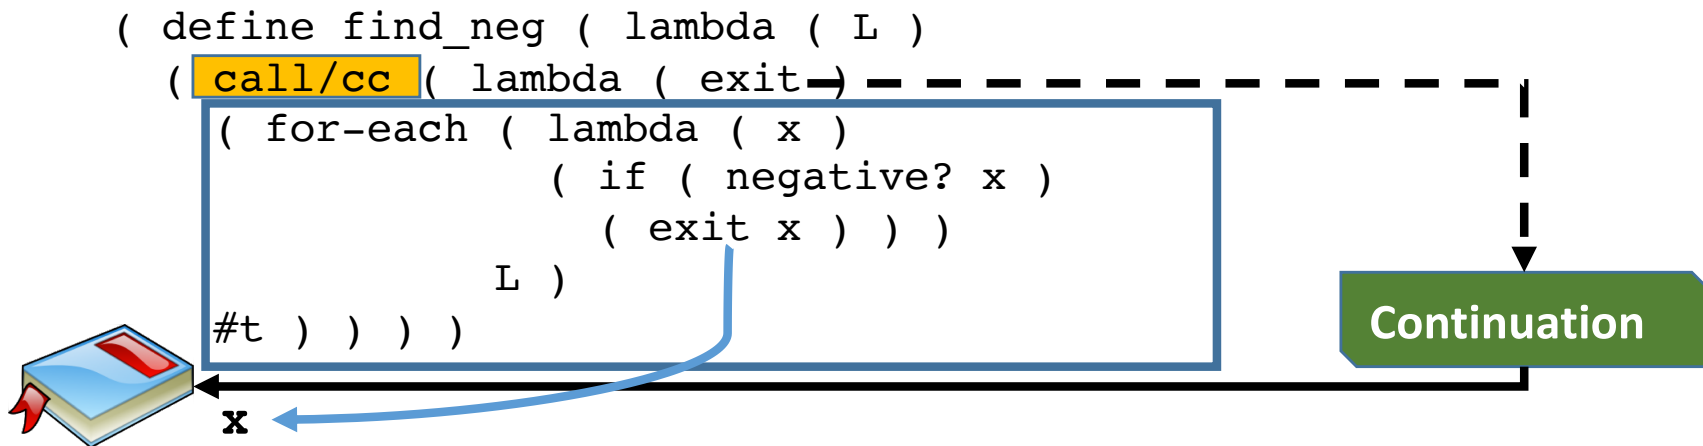


# Uses of Continuations

- Continuations can be used to support a variety of special purpose constructs in other languages
  - Escape procedures
  - Exception handling
  - Co-routines
  - Etc

# Escape Procedure with Continuation

- Simplest possible use: Escape procedure
  - If function  $f$  never makes use of the continuation, everything works as if (  $f$  ) was performed
  - If function  $f$  invokes the continuation, then program state is restored as if  $f$  was never called
- Example: Look for the first negative number in a list



- What happens when  
 ( find\_neg ' ( 54 0 37 -3 245 19 ) )  
=> -3

# Exception Handling with Continuations

- Suppose you want to sum a list of integers...

```
( define sum
  ( lambda ( lst )
    ( call/cc ( lambda ( exception )
      ( define r ( lambda ( L )
        ( cond
          ( ( null? L ) 0 )
          ( ( integer? ( car L ) )
            ( + ( car L ) ( r ( cdr L ) ) )
            ( #t ( exception #f ) ) ) ) )
      ( r lst ) ) ) ) ) )
```

Return to this  
continuation if an  
exception occurs

Empty list  
sum is 0

Else if next item  
is an integer

Else, next item is  
not an integer

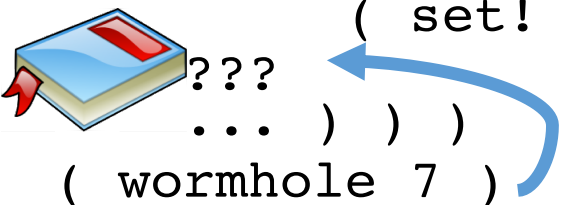
Call r with lst

```
( sum ( 1 2 3 4 ) )
=> 10
( sum ( 1 b 3 4 ) )
=> #f
```

# Gotos with Continuations

- But wait there is more!
- The continuation can be invoked anywhere!
- Even after we leave the scope where it was created
- What does this do?

```
( define wormhole #f )  
( define rabbit-hole  
  ( let ( ( x 42) )  
    ( call/cc ( lambda ( hole )  
                ( set! wormhole hole ) ) ) )  
    ... ) ) )  
( wormhole 7 )
```



The diagram illustrates the execution of the code. A blue book icon with a red bookmark is positioned to the left of the code. A blue arrow originates from the expression `( wormhole 7 )` and points to the line `( set! wormhole hole )` inside the lambda function, indicating that the continuation is invoked from a later point in the program.

- This is almost like a goto!
- Use at own risk!

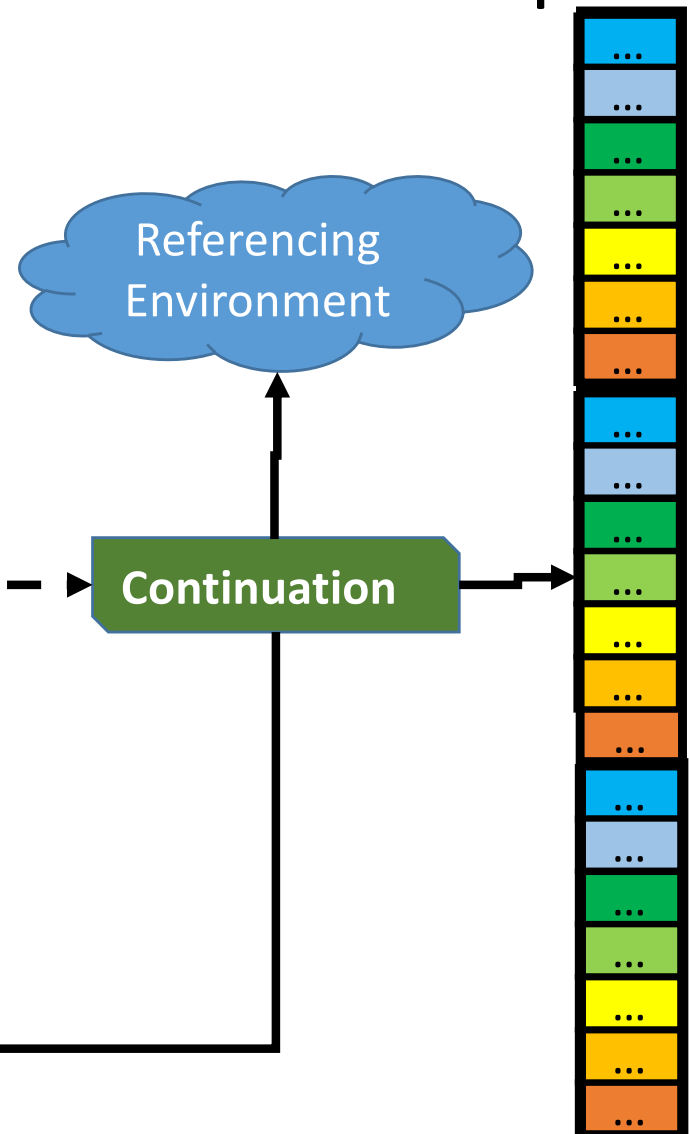
# Implementing Continuations

- To implement continuations, need to preserve:
  - Referencing environment : easy
    - Same as closures
  - Current program state: easy
    - Store current CPU state inside a continuation record
  - Stack content : Not so easy (it depends)
    - If continuations are only used within scope of creation, then pointer to stack frame is sufficient
    - If continuations used anywhere, need to make a copy of the entire stack! Why?

# Continuations Called within Scope

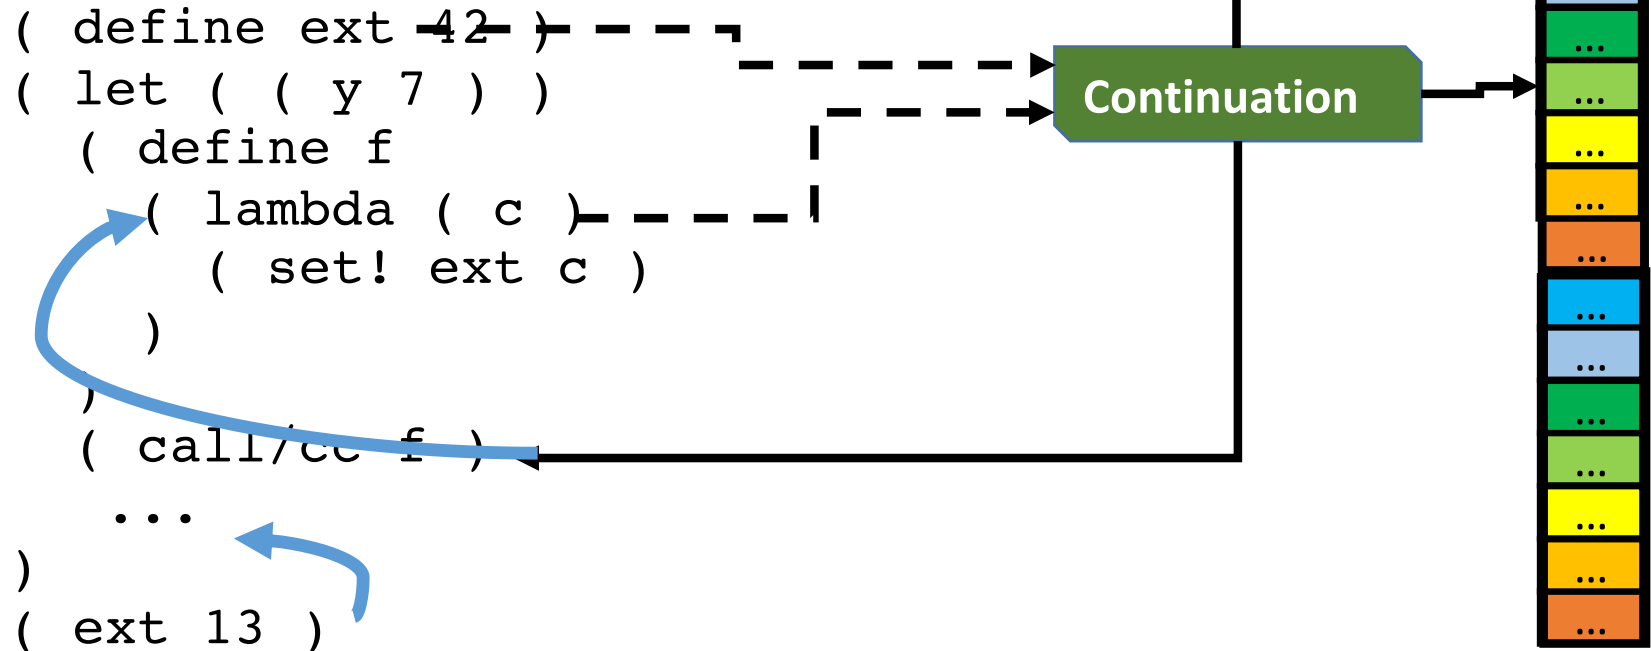
- If a continuation is called within the scope it was created, the stack frame present at its creation has not been destroyed

```
( define f  
  ( lambda ( c )  
    ...  
    ( c value )  
    ... )  
)  
...  
( call/cc f )
```



# The Challenge with Continuations Called Outside of Scope

- If a continuation is called outside the scope it was created, the stack frame present at its creation may be destroyed ☹️

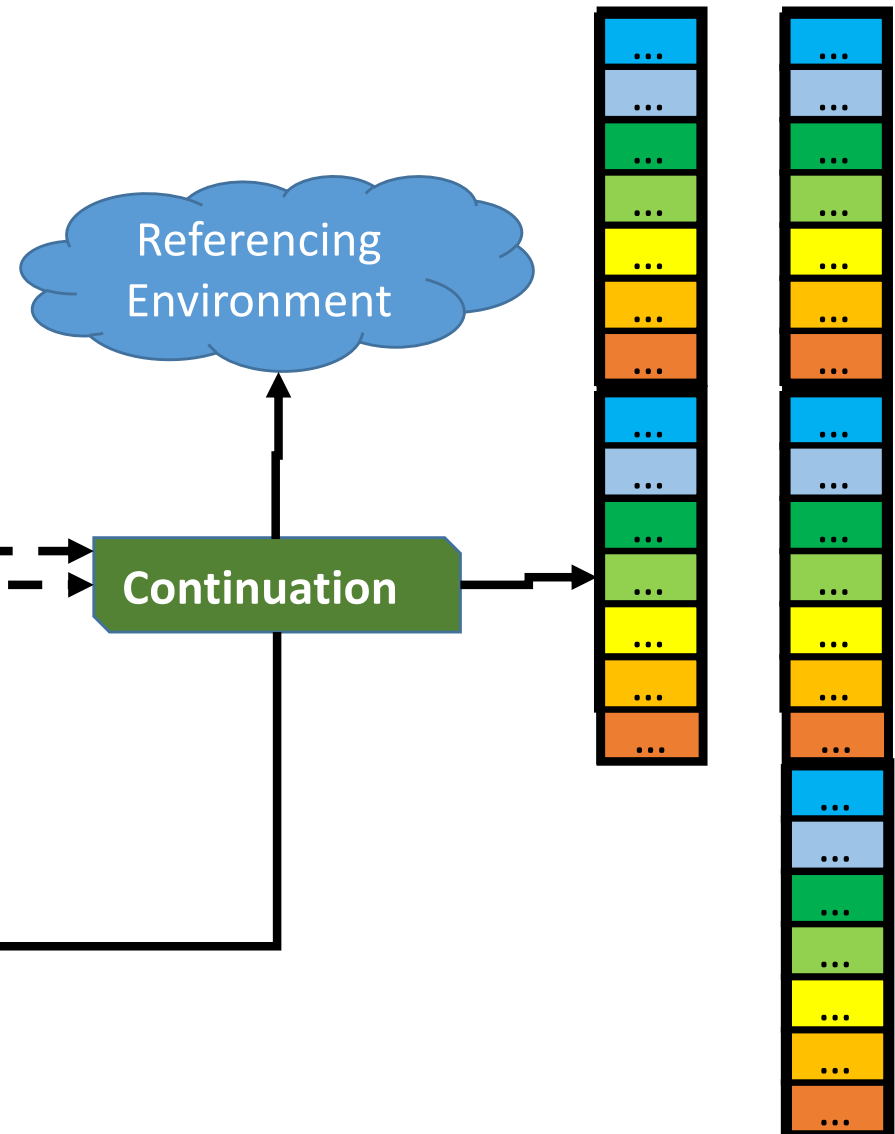




# Solution to Continuations Called Outside of Scope

- If a continuation is called outside the scope it was created, the stack frame present at its creation has to be duplicated

```
( define ext 42 )  
( let ( ( y 7 ) )  
  ( define f  
    ( lambda ( c )  
      ( set! ext c )  
    )  
  )  
  ( call/cc f )  
)  
( ext 13 )
```



# Recall setjmp()/longjmp()


- First invocation of `setjmp( )`
  - stores the current context in `env`
  - returns 0
- If no `longjmp(env, ...)` occurs
  - Branch terminates as usual
- If `longjmp(env, val)` occurs
  - It jumps directly into `setjmp( )`
- The `setjmp( )` returns again
  - And returns `val`
  - The else-branch is executed
- Recall, this is how we would implement exceptions in C

```
static jmp_buf env;
...
val = setjmp( env );
if( val == 0 ) {
    /* protected code */
    ...
    longjmp(env, val);
    ...
} else {
    /* handler */
}
...
```

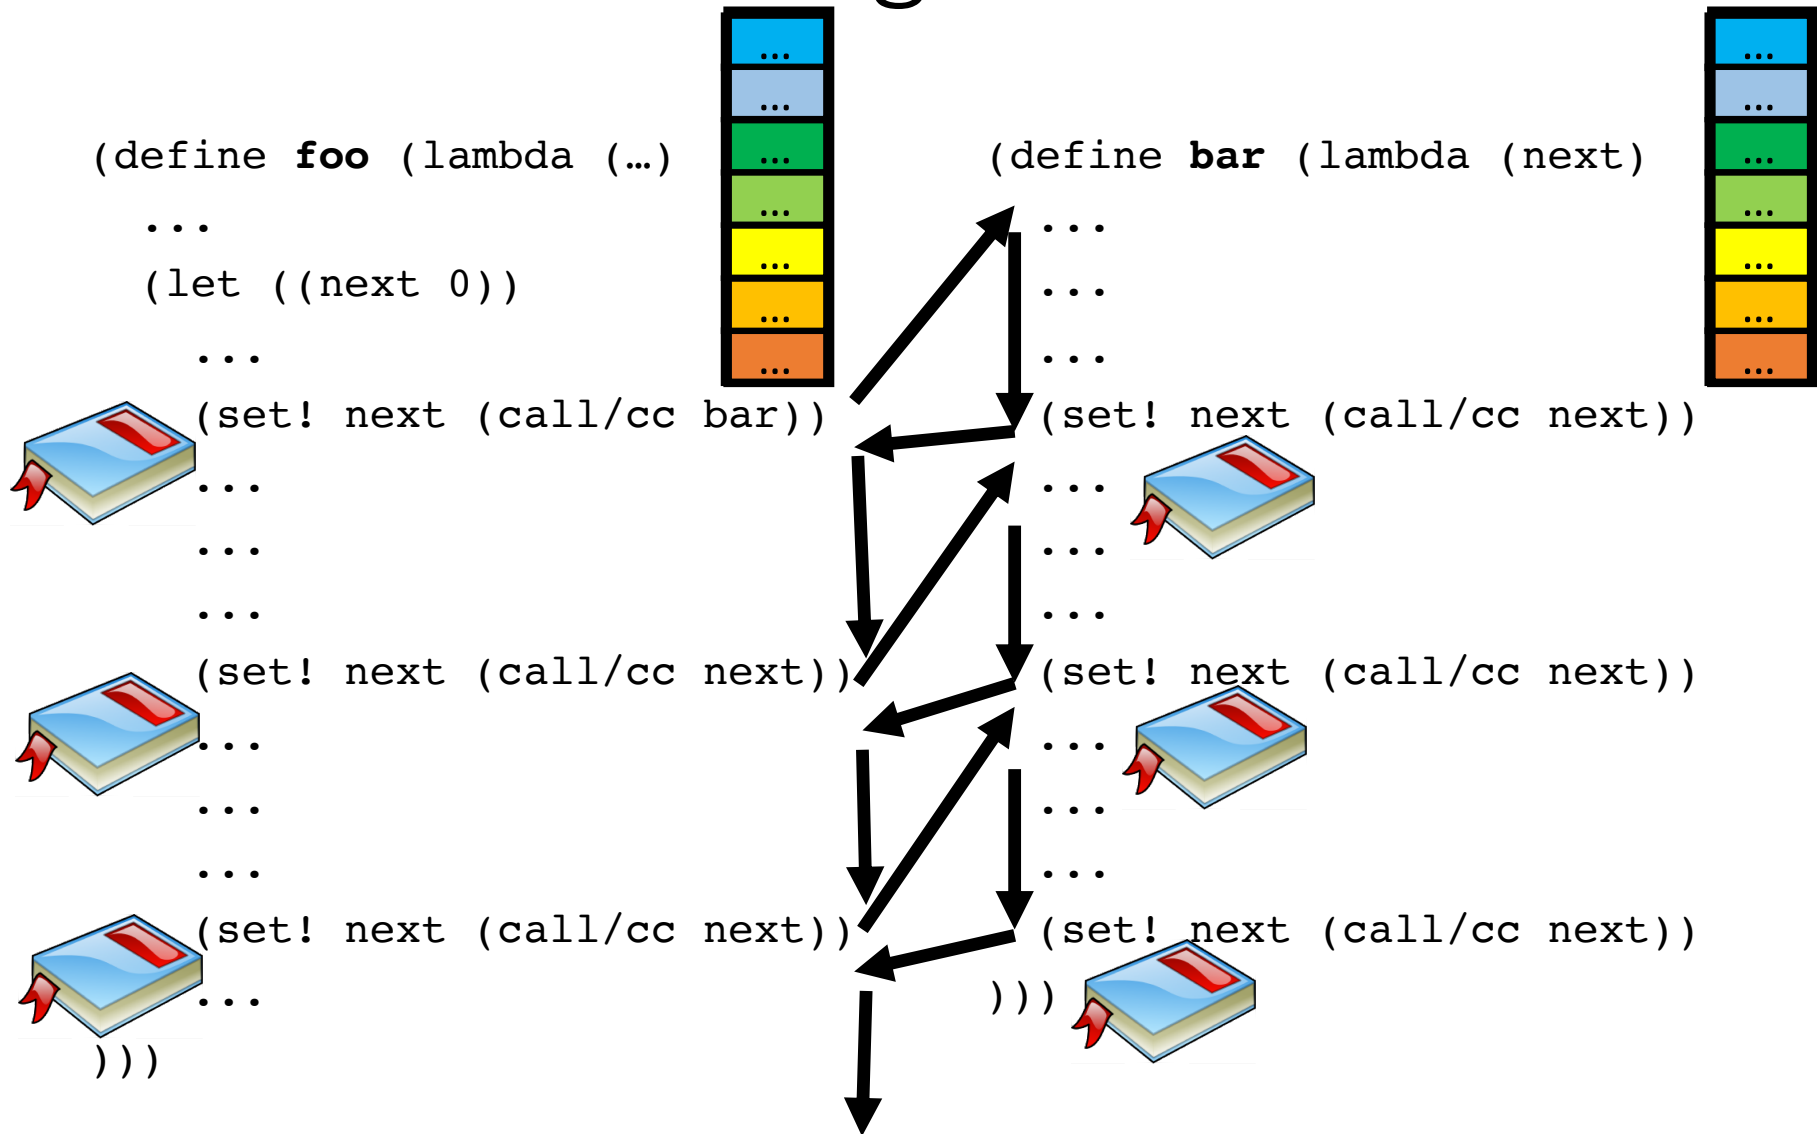
From the man page: *“The longjmp() routines may not be called after the routine which called the setjmp() routines returns.”*

# Coroutines

- Coroutines are separate threads of execution that yield control to each other
  - In contrast, real threads do not yield control
- Coroutines were commonly used to structure concurrent operations
- Useful for implementing
  - iterators
  - Generators
- Example:
  - Jumping back and forth from a generator
- Challenge: Need a separate stack for each coroutine.



# Coroutines using Continuations



# Discussion:

- Everything old is new again
    - Stephen King, and many others
  - Closures and continuations are not “new” concepts
  - They have been around for quite some time
  - They are now being rediscovered, implemented and used in modern languages because they are a useful way of specifying computation
  - These are not esoteric concepts that you will never use
  - Languages such as
    - Scala
    - Ruby
    - Actionscript
    - Python
    - Java
- Have these features or will have them!

# Coming Up Next

- Type Systems
- Data Types
- Arrays
- Garbage Collection